

1.

For this question, refer to the KnightMotives Automotive case study. KnightMotives management wants to upskill their engineering teams. They want each team to experiment with one specific technology at a time in a sandbox environment in Google Cloud. KnightMotives management also wants to prevent any experimental projects from going into production and creating shadow IT. They want you to design a solution that supports these training objectives and follows Google-recommended practices. What should you do?

- A. Create a request form where engineers can request a sandbox environment for a specific technology. Automate the creation of a project with only the relevant APIs enabled and lower the default API quota. Grant IAM roles related to this scope to the requesting team. Ensure the project is automatically deleted after a predefined amount of time.
- B. Create a request form where engineers can request a sandbox environment and automate the creation of a project in which the requesting engineer obtains full permissions. Within the details of the request form, highlight the expectation of keeping costs under control and the scope limited to the requested services.
- C. Create a separate billing account for all sandbox projects. Grant all engineers the role of Billing Account User on this billing account. Create a folder called Sandboxes in the Organization's Resource Manager. Grant all users the Project Creator role on this folder.
- D. Create a secondary Google Cloud organization which is to be used for R&D, separate from the production organization. Create a folder for each team under this organization. Create a project for each team member in the relevant folder. Grant each team member full access to their dedicated project. Grant each team manager the permission of Security Admin on the respective folder. Have them enable the relevant APIs in each team member's project.
- Answer

A

2.

For this question, refer to the KnightMotives Automotive case study. KnightMotives has developed and deployed a model on Vertex AI that can provide personalized recommendations in the new car configuration application. Customers will receive optional equipment recommendations that best suit their persona. Previous usage data from the car configuration application has been used for model training features. You know from past experience that customer behavior can change over time. For example, in times of economic certainty and rising stock markets, customers tend to purchase more expensive options. You want to detect when customer behavior gradually changes over time so you can adjust the model.

- A. Configure Model Monitoring, and select training-serving skew detection.
- B. Configure Model Monitoring, and select prediction drift detection.
- C. Configure Dataplex auto data quality on the prediction request data features using row-level rules.
- D. Configure Dataplex auto data quality on the prediction request data features using aggregate rules.
- Answer

B

3.

For this question, refer to the KnightMotives Automotive case study. KnightMotives has developed a new car configurator application to enhance both the dealer and customer experience. The new application will be deployed on Google Kubernetes Engine(GKE). KnightMotives wants you to replicate the deployment of the application across multiple locations to ensure this critical application is always available over the internet, even during regional outages. What should you do?

- A. Create multiple GKE clusters in different regions. Deploy part of the microservices of the app in different clusters. Configure a multi- cluster Cloud Service Mesh.

- B. Create multiple GKE clusters in different regions. Deploy the app on every cluster. Configure a multi- cluster Cloud Service Mesh.
- C. Create a fleet of GKE clusters in different regions. Deploy part of the microservices of the app in different clusters. Configure a multi- cluster Gateway.
- D. Create a fleet of GKE clusters in different regions. Deploy the app on every cluster. Configure a multi- cluster Gateway.
- Answer

D

4.

For this question, refer to the KnightMotives Automotive case study. As part of its development of fully autonomous driving vehicles, KnightMotives wants to analyze all vehicle sensor data during test drives. The analysis will enable KnightMotives to improve its software based on insights from this data. Different event types, such as parking, overtaking and navigating, need to be analyzed. Each test vehicle and event type has an ID. Different categories of sensors, such as cameras, radars, and ultrasonic beams, also need to be analyzed. During each autonomously-initiated event, data from multiple sensors will be captured and sent to Google Cloud where the data will be stored in Bigtable. During data analysis, you want to be able to retrieve all sensor data of occurrences of the same event type for a specific vehicle within a specific interval of time. You need to design a Bigtable schema that is optimized for read performance. What should you do?

- A. Use the sensor category, event ID, and timestamp(in that order) as the row key. Create a column family for each individual sensor and a column qualifier for each vehicle.
- B. Use the timestamp, vehicle ID, and event ID(in that order) as the row key. Create a column family per sensor category, and use a column qualifier for each individual sensor within its respective category.
- C. Use the vehicle ID and event ID(in that order) as the row key. Create a column family per sensor category, and C use a column qualifier for each individual sensor within its respective category. Utilize the timestamped versions of a cell to distinguish between different moments in time of similar events.

- D. Use the vehicle ID, event ID, and timestamp in that order) as the row key. Create a column family per sensor category, and use a column qualifier for each individual sensor within its respective category.
- Answer

D

5.

For this question, refer to the KnightMotives Automotive case study. KnightMotives has deployed their new Dealer Tools on Google Kubernetes Engine(GKE). You are supporting KnightMotives in the performance of the GKE workloads. Some of the Pods are Statefull. To ensure these Pods do not consume an unlimited amount of CPU, you have configured resource limits on the Deployments. You have identified that these Pods were allocated too much CPU and are actually consuming drastically less some of the time. You want to optimize resource usage and allow maximum flexibility in case of higher load. What should you do?

- A. Configure resource requests per Deployment. Set resource requests slightly above the typical CPU usage observed during monitoring.
- B. Configure resource requests per Deployment. Ensure the requests are equal to the current limits.
- C. Enable vertical Pod autoscaling on each Deployment. Ensure the autoscaling mode is set to Auto.
- D. Enable horizontal Pod autoscaling on each Deployment. Ensure the minimum number of Pods is lower than the currently configured amount in the Deployment.
- Answer

A

6.

For this question, refer to the KnightMotives Automotive case study. KnightMotives has established a dedicated Google Cloud Interconnect with 99.99% availability between its headquarters and two different metropolitan areas corresponding to the us-central1 and us-east1 Google Cloud regions. To minimize cost and latency between these workloads, you want to ensure all workloads in Google Cloud are deployed in these two regions as the VLAN attachment for the Cloud Interconnect. What should you do?

- A. Schedule an export of all assets via Asset Inventory into BigQuery. Schedule a daily Cloud Run function to query the export and send out an alert if resources are created in regions outside of the allowed list.
- B. Ensure all deployments are done through Infrastructure as Code using standardized Terraform modules. Configure the region variable of all resources with a default value corresponding to one of the allowed regions.
- C. Configure the ed regures arealite Restriction constraint organization policy at the organization level, and ensure
- D. Limit the Google Cloud VPC used by the Cloud Interconnect to only have subnets in the allowed regions.

- `Answer`

C

7.

For this question, refer to the KnightMotives Automotive case study. KnightMotives is managing supplier data and pricing in a central MySQL database at headquarters(HQ). Only personnel at HQ are allowed to change the data. Each local plant stores a copy of the data in their own MySQL database, often using a different database schema or version. Every night a batch job exports any product or price updates in XML format from the central database at HQ and stores the updated data on a central FTP server. Each local plant must download this XML file and update their local system with the new information. The local data kept by some plants has become inconsistent

with the source data due to XML parsing issues. HQ wants to easily verify that all changes are applied correctly at each plant.

- A. Create a Pub/Sub topic per supplier, and have HQ publish all changes related to the respective supplier in JSON format on that topic. Allow all plants to create a Pub/Sub Pull subscription to receive messages for their suppliers and update their databases.
- B. Create a Pub/Sub topic per supplier, and have HQ publish all changes related to the respective supplier in JSON format on that topic. Allow all plants to create a Pub/Sub Push subscription to receive messages for their suppliers and update their databases.
- C. Migrate the self-hosted MySQL database at HQ to Cloud SQL. Standardize the database schema across all plants, and configure the local databases as external read replicas in Cloud SQL.
- D. Migrate the self-hosted MySQL database at HQ to Cloud SQL. Configure Database Migration Service between Cloud SQL and the on-premises databases. When creating the migration job, choose continuous migration instead of one-time migration.
- Answer

C

8.

For this question, refer to the KnightMotives Automotive case study. KnightMotives wants to personalize the dealer experience for its customers and has decided to train its own AI models for personalized recommendations. The company will start collecting personally identifiable information(PII) from its customers to use as part of the models' training data. KnightMotives wants to ensure maximum security and compliance worldwide. You need to ensure the data is encrypted both at rest and during AI model training without impacting the models' accuracy. What should you do?

- A Store the training data in BigQuery using column-level encryption. Train the model using Confidential GKE Nodes.
- B. Store the training data in BigQuery using column-level encryption. Train the model on VertexAI notebooks using customer-managed encryption keys.

- C. Process all data with Sensitive Data Protection's de-identification service. Replace any PII with a random string before storing it. Train the model using Confidential GE Nodes.
- D. Process all data with Sensitive Data Protection's de-identification service. Replace any PII with a random string before storing it. Train the model on Vertex AI notebooks using customer-managed encryption keys.
- Answer

B

9.

For this question, refer to the KnightMotives Automotive case study. You are responsible for designing the network infrastructure architecture for KnightMotives's new environment on Google Cloud. You need to design the new VPC topology. You want to ensure a guaranteed bandwidth and low latency between the plants and Google Cloud resources. What should you do?

- A. Create a Standard Tier VPC, and ensure a subnet is available in the region closest to a plant. Establish Direct Peering between Google's Edge Network and the local plant.
- B. Create a Premium Tier VPC, and ensure a subnet is available in the region closest to a plant. Establish Direct Peering between Google's Edge Network and the local plant.
- C. Create a Premium Tier VPC, and ensure a subnet is available in the region closest to a plant. Establish a Cloud Interconnect between each subnet and the local plant.
- D. Create a Standard Tier VPC, and ensure a subnet is available in the region closest to a plant. Establish a Cloud Interconnect between each subnet and the local plant.
- Answer

C

10.

For this question, refer to the KnightMotives Automotive case study. As part of its cloud strategy,

KnightMotives wants to move its virtual machines (VMs) into Google Cloud. These VMs contain applications, database instances, and file servers. In the future, KnightMotives intends to modernize their landscape by adopting cloud native technologies on Google Cloud. During modernization, minimal latency and operational overhead between the old and new landscape will be crucial. You need to design the cloud architecture for the VM environment on Google Cloud while ensuring a future modernization track will meet the technical requirements. What should you do?

- A. Create a Shared Virtual Private Cloud (VPC), and migrate the VMs to Compute Engine in a new project. A During modernization, deploy each modernized workload in a dedicated Google Cloud project with its own VPC. Connect the two VPCs via an on-premises VPN.
- B. Create a Shared Virtual Private Cloud (VPC). Migrate the VMs to Compute Engine in a new project. During modernization, deploy each modernized workload in a dedicated Google Cloud project using the Shared VPC.
- C. Create a single private cloud on Google Cloud VMware Engine. Migrate the VMs using VMware HCX. During modernization, deploy the modernized workloads in a single Google Cloud project using a single VPC.
- D. Create a single private cloud on Google Cloud VMware Engine. Migrate the VMs using VMware HCX. During modernization, deploy the modernized workloads in a dedicated Google Cloud project with its own VPC.
- Answer

B

11.

You are planning to migrate your on-premises compute and SAP workloads to Google Cloud. You want to follow Google-recommended practices to quickly create a cost estimate for running these workloads in Google Cloud. What should you do?

- A. Use the Google Cloud pricing calculator, and input the estimated resource to generate a cost estimate.

- B. Engage with a Google Cloud partner to perform a comprehensive assessment and provide a customized cost estimate.
- C. Gather data about your current environment, and leverage Google Cloud Migration Center to generate a cost estimate.
- D. Leverage Cloud Asset Inventory to gather data and generate a cost estimate.
- Answer

C

12.

You are designing a central, automated infrastructure deployment process for your organization using Terraform and Cloud Build. The security team prohibits the use of long-lived, static service account keys in any CI/CD pipeline. Additionally, while developers can propose infrastructure changes for peer review, they must not have permissions to directly apply changes in the production project. You need to design a secure and automated workflow for applying Terraform changes that meets the security team's requirements and ensures proper governance. What should you do?

- A. Use service account impersonation in Cloud Build. Configure the pipeline to run terraform plan on pull requests, and require manual approval before running terraform apply.
- B. Configure the Cloud Build pipeline to use service account impersonation. Set up a trigger that automatically runs terraform apply when a pull request is merged.
- C. Configure the pipeline to only run terraform plan. After a pull request is approved, have an authorized developer run terraform apply from a secured workstation.
- D. Create a privileged service account and store its JSON key in Secret Manager. Configure the Cloud Build pipeline to fetch this key during execution to authenticate Terraform.
- Answer

A

13.

You need to build a continuous delivery pipeline for a containerized application in Google Cloud. You want to run all your tests in the pipeline to improve your application's quality. What should you do?

- A. 1. Run unit tests in the developer's local environment before committing and pushing the code to a central repository.
2. After the code is pushed, trigger Cloud Build to build the application container and deploy the container to a testing environment, and run integration tests.
3. If the integration tests are successful, deploy the container to your production environment and run acceptance tests.
- B. 1. Automatically run unit tests in a local environment once the code is changed. If all tests are successful, build a container.
2. Trigger Cloud Build to deploy the container to a testing environment, and run integration tests and acceptance tests.
3. If all tests are successful, tag the code as production ready. Trigger Cloud Build to build and deploy the container to the production environment.
- C. 1. After the developers push the code to a central repository, trigger Cloud Build to build the application container. Then run unit tests.
2. If unit tests are successful, deploy the container to a testing environment and run integration tests.
3. If the integration tests are successful, deploy the container to a production environment and run acceptance tests.
- D. 1. After the developers push the code to a central repository, trigger Cloud Build to run unit tests. If all unit tests are successful, build the application container and push it to a central registry.
2. Trigger Cloud Build to deploy the container to a testing environment and run integration tests and acceptance tests.

3. If all tests are successful, deploy the application to the production environment and run the smoke tests.

- Answer

C

14.

Your organization is a cloud-native shop which uses Google Kubernetes Engine(GKE) extensively to run its containerized workloads. You want to embrace infrastructure as code(IaC) principles to streamline and automate the provisioning, configuration, and management of Google Cloud resources and minimize operational cost. What should you do?

- A. Install Config Sync, and use YAML manifests in your GitHub repository to automate the provisioning and management of Google Cloud resources.
- B. Set up Config Controller, and use YAML manifests to define the desired state of your Google Cloud resources. Apply these YAML configurations to provision and manage the resources.
- C. Develop a custom Kubernetes operator within your GKE cluster to automate the execution of your IaC scripts for provisioning and managing Google Cloud resources.
- D. Install Config Connector, and use Python to define the desired state of your Google Cloud resources. Apply the Python scripts to provision and manage the resources.
- Answer

B

15.

You are deploying a critical application with a stateless, containerized frontend on Cloud Run and a Cloud SQL for PostgreSQL backend. The application experiences unpredictable

traffic spikes, and the business requires the ability to immediately roll back a failed deployment to the last known good state. You need to apply a deployment strategy that aligns with Site Reliability Engineering(SRE) principles for both the application code and the database schema updates, while meeting the business's requirements. What should you do?

- A. Package the database schema migration script within the container to be executed on every container startup before the application process begins.
- B. Separate CI/CD pipelines for database schema migrations from application deployments. When deploying a new Cloud Run revision, use gradual traffic split.
- C. Use a single CI/CD pipeline that first applies database schema changes and then deploys the new Cloud Run revision
- D. Configure the CI/CD pipeline to use the:latest container tag for deployments, with database schema changes applied manually as needed.
- Answer

B

16.

Your company sends all Google Cloud logs to Cloud Logging. Your security team wants to monitor the logs. You want to ensure that the security team can react quickly if an anomaly such as an unwanted firewall change or server breach is detected. You want to follow Google- recommended practices. What should you do?

- A. Export logs to a Pub/Sub topic, and trigger Cloud Function with the relevant log events.
- B. Schedule a cron job with Cloud Scheduler. The scheduled job queries the logs every minute for the relevant events.
- C. Export logs to a Cloud Storage bucket, and trigger Cloud Run with the relevant log events.
- D. Export logs to BigQuery, and trigger a query in BigQuery to process the log data for

the relevant events.

- Answer

A

17.

You are monitoring Google Kubernetes Engine(GKE) clusters in a Cloud Monitoring workspace. As a Site Reliability Engineer(SRE), you need to triage incidents quickly. What should you do?

- A. 1. Navigate the predefined dashboards in the Cloud Monitoring workspace.
2. Create custom metrics and install alerting software on a Compute Engine instance.
- B. 1. Navigate the predefined dashboards in the Cloud Monitoring workspace.
2. Add metrics and create alert policies.
- C. 1. Create a custom dashboard in the Cloud Monitoring workspace for each incident.
2. Add metrics and create alert policies.
- D. 1. Write a shell script that gathers metrics from GKE nodes, and publish these metrics to a Pub/Sub topic. 2. Export the data to BigQuery, and make a Data Studio dashboard.

- Answer

B

18.

You want to allow your operations team to store logs from all the production projects in your Organization, without including logs from other projects. All of the production projects are contained in a folder. You want to ensure that all logs for existing and new production projects are captured automatically. What should you do?

- A. 1. Create an aggregated export on the Organization resource.
- 2. Set the log sink to be a Cloud Storage bucket in an operations project.
- B. 1. Create an aggregated export on the Production folder.
- 2. Set the log sink to be a Cloud Storage bucket in an operations project.
- C. 1. Create log exports in the production projects.
- 2. Set the log sinks to be a Cloud Storage bucket in an operations project.
- D. 1. Create log exports in the production projects.
- 2. Set the log sinks to be BigQuery datasets in the production projects.
- 3. Grant IAM access to the operations team to run queries on the datasets.
- Answer

B

19.

You are managing an application deployed on Cloud Run, and you need to define a strategy for deploying new versions of the application. You want to evaluate the new code with a subset of production traffic to decide whether to proceed with the rollout. What should you do?

- A. 1. Deploy a new revision to Cloud Run with the new version.
- 2. Configure traffic percentage between revisions.
- B. 1. Deploy a new service to Cloud Run with the new version.
- 2. Add Cloud Load Balancing in front of both services.
- C. 1. In the Google Cloud console page for Cloud Run, set up continuous deployment using

Cloud Build for the development branch.

2. As part of the Cloud Build trigger, configure the substitution variable TRAFFIC_PERCENTAGE with the percentage of traffic you want directed to a new version.

- D. 1. In the Google Cloud console, configure Traffic Director with a new Service that points to the new version of the application on Cloud Run.

2. Configure Traffic Director to send a small percentage of traffic to the new version of the application.

- `Answer`

A

20.

Your company has an application deployed on Anthos clusters that is running multiple microservices. The cluster has both Anthos Service Mesh and Anthos Config Management configured. End users inform you that the application is responding very slowly. You want to identify the microservice that is causing the delay. What should you do?

- A Use the Service Mesh visualization in the Google Cloud console to inspect the telemetry between the microservices.
- B. Use Anthos Config Management to create a clusterselector selecting the relevant cluster. On the Google Cloud console page for Google Kubernetes Engine, view the workloads and filter on the cluster. Inspect the configurations of the filtered workloads.
- C. Reinstall Istio using the default Istio profile in order to collect request latency. Evaluate the telemetry between the microservices in the Google Cloud console.
- D. Use Anthos Config Management to create a NamespaceSelector selecting the relevant cluster namespace. On the Google Cloud console page for Google Kubernetes Engine, view the workloads and filter on the namespace. Inspect the configurations of the filtered workloads.
- Answer

A

21.

For this question, refer to the Altostrat Media case study. Altostrat stores a large library of media content, including sensitive interviews and documentaries, in Cloud Storage. They are concerned about the confidentiality of this content and want to protect it from unauthorized access. You need to implement a Google-recommended solution that is easy to integrate and provides Altostrat with control and auditability of the encryption keys. What should you do?

- A. Use customer-managed encryption keys(CMEK) for all Cloud Storage buckets storing sensitive media content. Implement fine-grained access control using IAM roles and groups to restrict access to sensitive buckets.
- B. Configure Cloud Storage to use server-side encryption with Google-managed encryption keys. Create a bucket policy to restrict access to only authorized Google groups and required service accounts.
- C. O Use Cloud Storage default encryption at rest. Implement fine-grained access control using IAM roles and groups to restrict access to sensitive buckets.
- D. Implement client-side encryption before uploading it to Cloud Storage. Store the encryption keys in a HashiCorp Vault instance deployed on Google Kubernetes Engine(GKE). Implement fine-grained access control to sensitive Cloud Storage buckets using IAM roles.
- Answer

A

22.

For this question, refer to the Altostrat Media case study. Altostrat is experiencing fluctuating computational demands for its batch processing jobs. These jobs are not time-critical and can tolerate occasional interruptions. You want to optimize cloud costs and address batch processing needs. What should you do?

- A. Set up standard VM instances.

- B. Use Cloud Run functions.
- C. Configure reserved VM instances.
- D. Deploy spot VM instances.
- Answer

D

23.

For this question, refer to the Altostrat Media case study. Altostrat's development team is using a microservices architecture for their application. You need to select the most suitable testing approach to ensure that individual microservices function correctly in isolation. What should you do?

- A. Execute integration testing.
- B. Perform end-to-end testing.
- C. Run unit testing.
- D. Use load testing.
- Answer

C

24.

For this question, refer to the Altostrat Media case study. Altostrat needs to analyze the performance of its media processing pipeline running on Java-based Cloud Run function. You need to select the most effective tool for the task. What should you do?

- A. Analyze the data via Cloud Profiler.
- B. Instrument the code to use Cloud Trace.
- C. Inspect data from Snapshot Debugger.

- D. Query logs in Cloud Logging.
- Answer

A

25.

For this question, refer to the Altostrat Media case study. Altostrat is concerned about sophisticated, multi- vector Distributed Denial of Service(DDoS) attacks targeting various layers of their infrastructure. DDoS attacks could potentially disrupt video streaming and cause financial losses. You need to mitigate this risk. What should you do?

- A. Set up VPC Service Controls to restrict access to sensitive resources and prevent data exfiltration.
- B. Configure Cloud Next Generation Firewall(NGFW) with custom rules to filter malicious traffic at the network
- C. Activate Security Command Center to monitor security posture and detect potential threats.
- D. Deploy Google Cloud Armor with pre-configured and custom rules for L3/L4 and L7 protection.
- Answer

D

26.

For this question, refer to the Altostrat Media case study. Altostrat is using Apigee for API management and wants to ensure their APIs are protected from overuse and abuse. You need to implement an Apigee feature to control the total number of API calls for cost management. What should you do?

- A. Activate XML threat protection.
- B. Configure Quota policies.

- C. Set up API key validation.
- D. Integrate OAuth 2.0 authorization.
- Answer

B

27.

Your team plans to use Vertex AI to develop and deploy machine learning models for various use cases for fraud detection, product recommendations, and customer churn prediction. You want to enhance the security posture of the Vertex AI and Workbench environment by restricting data exfiltration. What should you do?

- A. Enable VPC Flow Logs to monitor network traffic to and from Vertex AI services and to identify suspicious activity.
- B. Create a service perimeter and include aiplatform.googleapis.com and notebooks.googleapis.com as protected services.
- C. Create a service perimeter and include ml.googleapis.com and document.googleapis.com as protected services.
- D. Enable Private Google Access for the VPC network to allow Vertex AI services to access public Google services without traversing the public internet.
- Answer

B

28.

You have an application that uses Vertex AI Feature Store to manage and serve product features for real-time recommendations. You want to monitor the performance and health of the application. You need to understand the overall duration of a request. What should you do

- A. Monitor the Queries per second for your featurestore.
- B. Measure the Latency of your requests.

- C Observe the Request size in your featurestore.
- D Track the Online serving throughput of your requests.
- Answer

B

29.

A large healthcare provider's primary electronic health record(EHR) application runs on Compute Engine instances with a Cloud SQL for PostgreSQL database, all located in the us-west1 region. A new regulatory mandate requires you to implement and document a business continuity plan(BCP). This plan must ensure that the EHR application can be fully recovered and operational in a different geographical region with a recovery time objective(RTO) of two hours and a recovery point objective(RPO) of 15 minutes. You need to design a disaster recovery strategy that meets these strict BCP requirements. What should you do?

- A. Deploy active managed instance groups(MIGs) in both us-west1 and us-east1, fronted by a global external A HTTP(S) Load Balancer. For the database, use a cross-region read replica in us-east, and rely on load balancer health checks to automatically fail over all traffic during an outage.
- B. Use Terraform to define the application's compute infrastructure. During a disaster, configure the Cloud SQL database in us-west1 to use a cross-region read replica in us-east1, build the environment in us-east1, and promote the replica.
- C. Take daily snapshots of the Compute Engine disks and Cloud SQL database. Copy these snapshots to a Cloud Storage bucket in us-east. During a disaster, manually restore the virtual machines(VMs) and database from the latest snapshots.
- D. Deploy a regional MIG in us-west for high availability, and rely on the Google Cloud SLA to ensure the region remains online.
- Answer

A

30

A financial services company is decommissioning one of its on-premises data centers. As part of this initiative, the company needs to perform a one-time migration of 500 TB of historical transaction archives to a Cloud Storage bucket for long-term retention. The data center's internet egress is 1 Gbps, which is shared with critical business operations. You must complete the secure data transfer within a 60-day window to meet the decommissioning deadline. What should you do?

- A. Write a script that uses the `gcloud storage cp--parallel` command to upload the data in chunks over the public internet during off-peak hours.
- B. Provision a Partner Interconnect connection with a 10 Gbps capacity to accelerate the data transfer, and then use Storage Transfer Service.
- C. Order a Transfer Appliance, copy the data to the appliance using your high-speed local network, and ship it back to Google to upload the data into your Cloud Storage bucket.
- D. Use Storage Transfer Service to create an agent-based transfer job that moves the data from the on- premises file servers directly to the Cloud Storage bucket.
- Answer

C

31.

You are designing a new insurance claims processing application that will be deployed on Google Kubernetes Engine(GKE). Your company's compliance team requires a complete and non-repudiable audit trail for all administrative actions from day one. Your application must capture who deploys a new container image, who modifies the GKE cluster's configuration, and who interacts with running pods or Kubernetes secrets using `kubectl`. What should you do?

- A. Activate the Security Command Center Premium tier to analyze GKE logs and detect threats, vulnerabilities, and misconfigurations in real time.
- B. Enable GKE Audit Logging to send Kubernetes API server logs to Cloud Logging, and ensure Cloud Audit Logs are enabled for the project.

- C. Deploy a DaemonSet to every node in the GKE cluster that runs a logging agent to collect and forward all container logs to Cloud Logging.
- D. Enable Binary Authorization on the GKE cluster, and create a policy that requires all deployed container images to be signed by a trusted attester.
- Answer

B

32.

You are launching a data fabric service on Google Cloud in ten weeks and need to transfer 2 petabytes of historical data from your on- premises data center to Google Cloud. You want to securely transfer the data with low latency. What should you do?

- A. Request Transfer Appliances from Google Cloud. Transfer the data using the appliances.
- B. Export the data to portable disks and send the disks to Google Cloud for data import.
- C. Create a 1 Gbps VPN connection between your data center and Google Cloud. Use Storage Transfer Service to transfer the data.
- D. Create a 1 Gbps connection using direct interconnect. Use the Storage Transfer Service for on- premises data to transfer the data.
- Answer

A

33.

You are designing the storage configuration for a critical Cloud Storage bucket. This bucket will store trade confirmation documents for a financial services company. These documents are immutable and serve as a legal record. Your design must meet two primary data protection requirements: The system must allow for the rapid recovery of any document that is accidentally deleted or maliciously overwritten. Once a document is written, it must be impossible for

anyone, including administrators, to delete or modify it for seven years. What should you do?

- A. Enable Object Versioning on the bucket. Additionally, configure the bucket with a Bucket Lock and a retention A policy of seven years.
- B. Configure an Object Lifecycle Management rule on the bucket to ensure all objects are retained for at least seven years before they can be deleted.
- C. Grant only the roles/storage.objectCreator IAM role to the application's service account. Remove all other admin permissions from the bucket.
- D. Use Storage Transfer Service to create a daily replication job that synchronizes the bucket to a separate archive bucket in another region.
- Answer

A

34.

Your company has hired an external auditing firm to perform a compliance audit. Your company's governance policy requires that external auditors be managed in a single Google Group that is granted temporary, read-only access to a Cloud Storage bucket named audit-evidence- bucket. Access must be traceable to the individual auditor's identity and be active only for the duration of the audit engagement, which runs the entire month of October. You need a secure access control strategy that avoids administrative overhead and complies with your company's governance policy. What should you do?

- A. Use Cloud Scheduler to run a Cloud Run functions script that adds the IAM binding of roles/ storage. objectviewer to the Google Group on October 1 and another that removes the IAM binding on November 1,
- B. Apply an IAM policy binding that grants the roles/storage.objectViewer role to the Google Group. Configure this binding with a time- based IAM Condition that automatically grants access from October 1 to November 1,
- C. Create a service account, and grant it the roles/storage.objectviewer role on the bucket. Generate and share Signed URLs for each object in the bucket with an expiration date of November 1.

- D. Use Workforce Identity Federation to map the auditors' group to the Google Group. Bind the roles/ storage.objectViewer role to this Google Group. Configure a 1-month session duration on the provider.
- Answer

B

35.

You are using a GitHub repository for your application's source code. You want to set up an efficient and secure continuous deployment process to automatically build and deploy the application to Cloud Run whenever a pull request is merged. What should you do?

- A. Create a workflow using GitHub Actions to build and deploy the application to Cloud Run once a pull request is merged. The workflow will use a service account key checked in with your source code for deployment permission.
- B. Create a Github webhook trigger in Cloud Build. Once a pull request is merged trigger Cloud Build to build a container image and save it in Artifact Registry. Use Config Sync to deploy the application to Cloud Run.
- C. Connect your repository using the Cloud Build GitHub app. Create a trigger in Cloud Build. Once a pull request is merged, trigger Cloud Build to build and deploy the application to Cloud Run.
- D. Create a GitHub Enterprise trigger in Cloud Build. Once a pull request is merged, trigger Cloud Build to build and deploy the application to Cloud Run. Save the deployment credential to Secret Manager.
- Answer

C

36.

You are deploying a highly confidential data processing workload on Google Cloud. Your company's compliance framework mandates that cryptographic keys used for encrypting data at

rest must be generated and stored exclusively within a validated Hardware Security Module (HSM). You want to use a fully integrated Google Cloud managed service to handle the lifecycle and usage of these keys. What should you do?

- A. Create a new key in Cloud Key Management Service (Cloud KMS) with the HSM protection level.
- B. Configure Cloud External Key Manager(Cloud EKM) to connect to your on-premises HSM.
- C. Use Customer-Supplied Encryption Keys(CSEK) by providing your on-premises generated key with each API request.
- D. Import your on-premises HSM key material into a Cloud KMS key with the SOFTWARE protection level.
- Answer

A

37.

Your team is running applications on a Google Kubernetes Engine(GKE) cluster with a private endpoint. You've set up a Cloud Deploy pipeline, but deployments to the KE cluster are failing. You need to resolve the issue. What should you do?

- A. Create a custom target in Cloud Deploy. Update the deploy pipeline to use the custom target for the application deployment.
- B. Create a Cloud Build private pool in the default VPC. Use Cloud Build to deploy the applications to the GKE cluster.
- C. Use VPC Service Controls with Cloud Build. Update the Cloud pipeline to use Cloud Build as its execution environment.
- D. Create a Cloud Build private pool that is peered with the same VPC network as your GKE cluster. Update the Cloud Deploy pipeline to use this private pool as its execution environment.
- Answer

D

38.

Your ecommerce platform uses a regional Cloud SQL for PostgreSQL database to store critical order information. The business requests a recovery time objective(RTO) of less than 10 minutes and a recovery point objective(RPO) of 2 minutes to ensure business continuity in the event of a full regional outage. You need to design a disaster recovery strategy for the Cloud SQL database that meets the business's strict RTO and RPO requirements. Your design must avoid operational complexity. What should you do?

- A. Schedule hourly automated backups of the Cloud SQL instance to a multi-regional Cloud Storage bucket. In a disaster, restore the latest backup to a new instance in a different region.
- B. Configure the Cloud SQL instance with a cross-region read replica. In a disaster, promote the read replica to a standalone, primary instance.
- C. Configure the primary instance for high availability(HA). In the event of a regional outage, trigger a failover to the standby instance.
- D. Use Database Migration Service to continuously replicate the database to another instance in a different region. In a disaster, redirect application traffic to the replica.
- Answer

B

39.

Your team is developing a new application using a microservices architecture on Kubernetes Engine. As part of the development lifecycle: Any code change that is pushed to the remote develop branch on your GitHub repository should be built and tested automatically. When the builds and tests are successful, the relevant microservice will be deployed automatically in the development environment. You want to ensure that all code deployed in the development environment follows this process. What should you do?

- A. 1. Have each developer install a pre-commit hook on their workstation that tests the code and builds the container when committing on the development branch.
- 2. After a successful commit, have the developer deploy the newly built container image on the development cluster.
- B. 1. Install a post-commit hook on the remote git repository that tests the code and builds the container when code is pushed to the development branch.
- 2. After a successful commit, have the developer deploy the newly built container image on the development cluster.
- C. 1. Create a Cloud Build trigger based on the development branch that tests the code, builds the container, and stores it in Artifact Registry.
- 2. Create a deployment pipeline that watches for new images and deploys the new image on the development cluster.
- 3. Ensure only the deployment tool has access to deploy new versions.
- D. 1. Create a Cloud Build trigger based on the development branch to build a new container image and store it in Artifact Registry.
- 2. Rely on Vulnerability Scanning to ensure the code tests succeed.
- 3. As the final step of the Cloud Build process, deploy the new container image on the development cluster.
- 4. Ensure only Cloud Build has access to deploy new versions.
- 5. Answer
- 6. C

40.

Your company uses the Firewall Insights feature in the Google Network Intelligence Center. You have several firewall rules applied to Compute Engine instances. You need to evaluate the efficiency of the applied firewall ruleset. When you bring up the Firewall Insights page in the Google Cloud console, you notice that there are no log rows to display. What should you do to

troubleshoot the issue?

- A. Enable Firewall Rules Logging for the firewall rules you want to monitor.
- B. Install the Google Cloud SDK, and verify that there are no Firewall logs in the command line output.
- C. Enable Virtual Private Cloud(VPC) flow logging.
- D. Verify that your user account is assigned the compute networkAdmin Identity and Access Management (IAM) role.
- Answer

A

41.

Your company wants to optimize Google Cloud costs for their development and staging environments. These environments are workstations used by developers Monday through Friday, 9:00 AM to 6:00 PM local time. Currently, the environments run on a fleet of n1-standard-4 Compute Engine instances that operate 24/7, leading to a high monthly cost for resources that are idle more than 70% of the time. You need to implement a solution that significantly reduces the monthly cost of these non-production environments without impacting the development team's productivity during work hours. What should you do?

- A. Purchase three-year committed use discounts(CUDs) for the existing n1-standard-4 instances.
- B. Re-architect the environments to run on a regional managed instance group(MIG) with autoscaling enabled.
- C. Schedule the virtual machines to start and stop to match your team's work schedule.
- D. Upgrade all instances to the N2 machine series.
- Answer

C

42.

A retail company's most critical application is its online payment processing system. The business has a requirement that the system must be able to survive a complete zonal outage while minimizing cost. You need a design solution that can handle a zonal failure. What should you do?

- A. Deploy the application in an active-active configuration using managed instance groups(MIGs) in two different (A regions, fronted by a global external HTTP(S) Load Balancer and backed by a multi-regional database like Spanner.
- B. Deploy the application on a regional MIG to provide high availability across multiple zones in the primary region.
- C. Configure the regional MIG to use only Spot VMs to aggressively minimize operational costs while maintaining high availability.
- D. Deploy the application on Compute Engine instances across multiple regions, and rely on daily snapshots for recovery to achieve the lowest possible cost.
- Answer

A

43.

A global media company is launching a new web application. The application backend is hosted on Compute Engine in us-central1 and serves both static assets(images, CSS, and Javascript) and dynamic, user-specific content from a Cloud SQL database in the same region. Early user feedback from Europe and Asia indicates significant page load delays due to slow loading static content. You need to design a solution that minimizes latency for all global users accessing the static content. What should you do? Choose 2 answers

- A. Deploy the application frontend service to Compute Engine managed instance groups in regions in Europe and Asia. Use a global external HTTP(S) Load Balancer to route user traffic to the nearest region.
- B. Create Cloud SQL read replicas in regions in Europe and Asia, and direct all database read traffic from those continents to their local replica.

- C. Vertically scale the Compute Engine instances in us-central by increasing their machine size.
- D. Use a regional external Network Load Balancer in us-central1 to better distribute the incoming global traffic.
- E. Enable Cloud CDN for the backend service that serves the static assets, and configure it as part of a global external HTTP(S) Load Balancer.
- Answer

AE

44.

You are developing a deep learning model that requires high-performance access to large volumes of media data currently stored in Cloud Storage. Model training will be executed on multiple VM instances with GPUs attached, but your application must interact with the data as if it were on the local file system. You need to minimize complexity and cost. What should you do?

- A. Copy the data from Cloud Storage to Filestore, and then mount the Filestore volume as a local file system on your VM instances.
- B. Create a shared persistent disk, attach the disk to your VM instances, and load data from the Cloud Storage bucket.
- C. Use the gcsfuse command line tool to mount the Cloud Storage bucket as a local file system, and perform read/write operations in your bucket using standard file system semantics.
- D. Use the gsutil command line tool to download the data to your VM instances.
- Answer

C

45.

Your company is expanding its AI-powered operations nationwide and has chosen accelerator-

based compute for the AI workloads. The batch image processing workloads are not time-sensitive and can tolerate interruptions. You need to rapidly deploy cost-effective accelerator nodes for these batch tasks, ensuring rapid deployment and data persistence when necessary. What should you do?

- A. Deploy Cloud Run functions with ephemeral local SSD.
- B. Deploy standard VMs with configured accelerators and attached persistent disks.
- C. Deploy spot VMs with attached persistent disks and implement checkpoint mechanisms.
- D. Deploy spot VMs with local SSD to reduce time for bursty workloads.
- Answer

C

46.

A large, multinational corporation is migrating to Google Cloud. The company has several distinct business units: Finance, Marketing, and Research and Development(R&D). The central security team has mandated governance requirements for each business unit: Finance: Must be restricted to deploying resources only in specific, compliant regions(us-central1 and europe-west2). Access to their projects must be tightly controlled by a dedicated finance-admins group. Marketing: Needs separate environments for production and development, with different teams managing each environment. R&D: Requires maximum flexibility to experiment with new services but must be completely isolated to prevent any impact on production systems. Global Auditing: A central compliance team requires read-only access to view all resources across the entire company for auditing purposes. You need to design a resource hierarchy that enforces these security policies at scale according to the Google Cloud Well-Architected Framework while providing the correct level of autonomy for each business unit. What should you do?

- A. Create a folder for each department under the root Organization node. Apply the resource location Organization Policy on the Finance folder. Within the Marketing folder, create separate projects for mktg-prod and mktg-dev. Grant the compliance team the roles/viewer role at the Organization level

- B. Place all projects directly under the Organization node. Use network tags and service accounts to enforce security boundaries between the different department workloads. Apply the resource location Organization Policy on the Finance project.
- C. Create separate google cloud Organizations for each department(Finance, Marketing, and R&D).Grant the compliance team the roles/viewer role for each organization.
- D. Create a single project for each department. Apply the resource location policy directly to the Finance project.Grant the compliance team the roles/browser role on each project individually.
- Answer

A

47.

Your organization is going to migrate applications to Kubernetes and use managed cloud services to deploy applications. Your team is new to Kubernetes and wants to quickly onboard engineers. You want to reduce operational overhead, so the engineering team can focus on developing consumer requirements instead of maintaining the infrastructure. What should you do?

- A Package your application into a Docker image, and deploy it to Kubernetes on Compute Engine.
- B. Leverage Cloud Build to create a container image, and deploy it automatically to Kubernetes on Compute Engine.
- C. Assess application and dependencies for containerization. Develop a migration strategy for deployment to GKE in Standard mode.
- D. Assess application and dependencies for containerization. Develop a migration strategy for deployment to GKE in Autopilot mode.
- Answer

D

48.

Your company runs a critical, revenue-generating ecommerce application that is served by a regional managed instance group(MIG) behind an external HTTP(S) Load Balancer. The operations team is currently overwhelmed with low-priority notifications and is starting to ignore alerts. Your team's service level objective(SLO) is to maintain 99% availability, which is measured by the ratio of successful requests (2xx status codes) to total requests. You want to minimize noise from non-critical events and ensure that the team is only notified of issues that are actionable and threaten the SLO. What should you do?

- A. Focus on cause-based alerts, creating alerting policies with thresholds for the Compute Engine instances, including CPU utilization, memory usage, disk I/O, and network traffic.
- B. Implement an error budget policy based on the availability of the SLO. Create a "page" alert that triggers only when the rate of burn of the error budget predicts a full exhaustion within the next 24 hours.
- C. Configure alerts based on predictive metrics. Use the instance count of the MIG as the primary metric to trigger an alert.
- D. Create log-based alerts for only the WARN and ERROR log entries generated by the application to ensure that no potential issue is missed.
- Answer

B

49.

You manage a highly distributed, hybrid- and multi-cloud IT environment, and your developers rely heavily on Prometheus for their workflows. You need a cloud- based, highly scalable, low-maintenance enterprise solution that supports Prometheus Query Language(PromQL) queries, quick metric viewing, and efficient issue diagnosis. What should you do?

- A. Set up Cloud Monitoring as a single pane of glass across multi-cloud environments.
- B. Deploy a Prometheus operator in your existing Kubernetes and Serverless setup across multi-cloud environments.
- C. Build a SaaS-based, Prometheus-compatible solution to display metrics for each

cloud in a customizable way.

- D. Enable Google Cloud Managed Service for Prometheus to monitor and alert on your workloads at scale.
- Answer

D

50.

Your organization uses Google Kubernetes Engine (GKE) and Amazon Elastic Kubernetes Service (EKS) to manage a complex Kubernetes environment across multiple cloud providers. You need to deploy a solution that streamlines configuration management, enforces security policies, and ensures consistent application deployment across all of the environments. You want to follow Google-recommended practices. What should you do?

- A. Deploy Kustomize for configuration customization, Config Sync with multiple Git repositories, and a script to enforce security policies.
- B. Leverage Argo CD for GitOps-based continuous delivery and Open Policy Agent(OPA) for policy enforcement, and develop a controller for multi-cluster configuration management.
- C. Utilize Config Sync as part of GKE to synchronize configurations from a centralized repository, and utilize Policy Controller to enforce policies using OPA Gatekeeper.
- D. Deploy Crossplane for managing cloud resources as Kubernetes objects, FluxCD for GitOps-based configuration synchronization, and Kyverno for policy enforcement.
- Answer

C

51.

Your company is rapidly deploying containerized microservices on Google Kubernetes Engine(GKE) using a robust CI/CD pipeline. Security is a top priority, and you need to implement a

comprehensive and efficient strategy to prevent container image vulnerabilities from reaching your GKE production environment. What should you do? Choose 2 answers

- A. Enable Shielded GKE Nodes on the production cluster to automatically block the execution of container images with known vulnerabilities.
- B. Incorporate vulnerability scanning before building container images, and use Google-maintained base images for your container deployments.
- C. Enable Artifact Analysis for the container images, and stop deployment if critical vulnerabilities are found.
- D. Use a custom security policy within your container image that restricts access to specific network ports and resources.
- E. Review the security reports generated by Artifact Analysis for each container image before deployment to GKE.
- Answer

BC

52.

Your organization is implementing a new cloud-native application on Google Cloud and needs to ensure compliance with the ISO/IEC 27001 framework. You want to leverage Google Cloud's security reports and documentation to support your ISO/IEC 27001 audit process. What should you do?

- A. Utilize the Cloud Audit Logs service for accessing and requesting the ISO/IEC 27001 reports.
- B. Review the Compliance Reports Manager for information about ISO/IEC 27001 compliance and related documentation on obtaining reports through your Google Cloud account.
- C. Engage an independent auditor to conduct an ISO/IEC 27001 audit of your organization's Google Cloud implementation.
- D. Download the ISO/IEC 27001 report for Google Cloud through internet search.

- Answer

B

53.

Your company is a global financial services provider that processes and analyzes a high volume of credit card transactions in real time for fraud detection. Your analytics team must run complex batch queries on the same transaction data for daily reporting. You need to design a data processing solution that can handle both real-time and batch processing of the transaction data while minimizing operational overhead and infrastructure management. What should you do?

Choose 2 answers

- A Use a Dataproc cluster for both the streaming and batch workloads.
- B. Use Dataprep to ingest the transactions.
- C. Use BigQuery for the batch analytics reports.
- D. Use Dataflow to process the streaming data.
- E. Use Firestore to store and analyze the transaction data.
- Answer

C D

54.

Your company uses a custom-built application running on a Compute Engine virtual machine(VM). This application processes real-time sales data and writes it to a zonal Persistent Disk. A recent internal audit requires that you implement a backup and recovery plan to protect against zonal failures. Your company has a strict policy that all backup data must be retained for at least 90 days and stored in a separate project with limited access. You need to implement a fully automated backup solution that meets these requirements with minimal operational overhead. What should you do?

- A. Write a script to create daily backups of the Persistent Disk. Copy the backups to a

different zone and apply a label to each snapshot to indicate the deletion date.

- B. Create a snapshot schedule to automatically create Persistent Disk snapshots and use a script to move and store them in a multi- regional Cloud Storage bucket.
- C. Use gcloud commands to create snapshots of the Persistent Disk. Store the snapshots in a regional Cloud Storage bucket and configure a lifecycle rule to delete objects older than 90 days.
- D. Use the Backup and Disaster Recovery(DR) service to create a backup plan. Configure the backup plan to take daily snapshots and store them in a backup vault with a 90-day retention policy.
- Answer

B

55.

You are architecting a new application feature for a healthcare provider based in Qatar. The feature needs to summarize sensitive patient notes that are submitted by clinicians. A critical requirement is that the content of these patient notes must never be processed outside of Qatar's borders. You want to use a powerful, pre-trained generative model for the summarization task while strictly adhering to the data residency constraint. What should you do?

- A. Use Vertex AI Model Garden to select a Gemma model. Deploy this model to a Vertex AI Endpoint within a Google Cloud region located in Qatar.
- B. Gather a large, anonymized dataset of medical notes. Use Vertex AI Training to train a custom summarization model from scratch, deploying it in a Qatar region.
- C. Use the Cloud Natural Language API to analyze the text and configure it to generate a summary of the patient notes.
- D. Write application code that sends the patient notes explicitly to the Gemini API endpoint in Qatar for summarization. Protect the API by VPC-Service Controls.
- Answer

A

56.

(33t) Q56. Your ecommerce web application includes many products with pictures and videos. You need to improve the application's page load speed and reduce the latency for customer requests. What should you do? Choose 2 answers

- A. Deploy your application to a GKE Autopilot cluster instead of a regular GKE cluster.
- B. Deploy your application to Cloud Run and enable concurrency.
- C. Change the database from Bigtable to Spanner.
- D. Use Memorystore for Redis.
- E. Implement CDN with the application's external HTTPS load balancer.
- Answer

BE

57.

You need to build and deploy a containerized web application to Google Cloud. The application is very write-heavy and requires a relational database as its data store. The application needs to be highly available in multiple cloud regions. You want to minimize operational overhead while following Google-recommended practices. What should you do?

- A. Deploy the application to a GKE standard cluster behind a global HTTPS load balancer. Use Firestore as the database with multi-region locations enabled.
- B. Deploy the application to GKE clusters in multiple regions and expose the application endpoint with a regional load balancer. Use Cloud SQL as the database.
- C. Deploy the application to a GKE standard cluster with a global ingress to expose the application endpoint. Use Bigtable as the database.
- D. Deploy application to Cloud Run in multiple regions behind a global HTTPS load balancer. Use Spanner as the database.
- Answer

D

58.

Your machine learning(ML) engineers use self-hosted Jupyter notebooks for tasks such as data preparation, model training, and fine-tuning. The operations team then deploys these models in various environments. You want to provide maximum flexibility for ML engineers, promote collaboration with a common toolset, and leverage Google Cloud's scalability, while following Google-recommended practices. What should you do?

- A. Use AutoML for machine learning and Cloud Deploy for model deployment.
- B. Use Colab Enterprise for machine learning and DevOps for model deployment.
- C. Use TensorFlow for machine learning and Cloud Deploy for model deployment.
- D. Use Vertex AI for machine learning and machine learning operations(MLOps) for model deployment.
- Answer

D

59.

Your employer is a financial services company that recently acquired a popular fintech startup. The startup's core application is a monolithic Python application running on a managed instance group of Compute Engine virtual machines with a single, large PostgreSQL database. Your development team struggles with slow deployment cycles, and the monolithic design of the startup's core application makes it difficult to integrate new, ML-powered fraud detection models. You need a long-term strategy that improves developer agility and positions the company to leverage Google Cloud's advanced data and AI capabilities for future innovations. What should you do?

- A. Migrate the PostgreSQL database to Cloud SQL for PostgreSQL. Replicate the data into BigQuery using Datastream, and then train and deploy the fraud detection models directly within BigQuery using BigQuery ML.

- B. Propose a phased, event-driven migration to a microservices architecture. Use Pub/Sub for asynchronous communication and deploy the fraud models on Vertex AI endpoints.
- C. Deploy the ML fraud detection model to a Vertex AI endpoint. Create a REST API for the model and modify the monolithic Python application to call this endpoint for real-time fraud analysis.
- D. Containerize the application, deploy it to Google Kubernetes Engine(GKE), and migrate the PostgreSQL database to Cloud SQL for PostgreSQL.
- Answer

B

60.

You are deploying a new three-tier application to Compute Engine instances within a single Virtual Private Cloud(VPC). The architecture is segmented into three subnets: a web tier subnet, an application tier subnet, and a database tier subnet. The web tier must only receive traffic from an external load balancer. The application tier must only receive traffic from the web tier. The database tier must only receive traffic from the application tier. You need to enforce strict traffic flow control and want to follow Google-recommended practices. What should you do?

- A. Create a single VPC firewall rule with a high priority that allows ingress traffic on all ports between the IP address ranges of the three subnets.
- B. Set up VPC Network Peering between the web tier and application tier subnets, and another peering connection between the application tier and database tier subnets.
- C. Use a combination of network tags and service accounts. Apply a unique network tag and a dedicated service account to the instances in each tier. Then create specific firewall rules that allow ingress traffic based on the source service account or tag of the upstream tier.
- D. Configure Cloud NAT for each subnet, and create Google Cloud Armor policies to filter traffic between the tiers based on their IP address ranges.
- Answer

C

61.

For this question, refer to the EHR Healthcare case study. Due to the growing partnerships with new insurance providers, EHR Healthcare needs a solution to streamline the creation of secure Google Cloud environments for each provider. You need to automate the provisioning and management of these secure environments, ensure that each provider's environment is consistently configured and managed, and provide a consistent user experience across all providers. What should you do?

- A. Utilize Infrastructure as Code(IaC) tools, such as Terraform, to define and manage infrastructure configurations.
- B. Leverage Google Kubernetes Engine(GKE) for workload deployments. Use cluster autoscaling and node auto-provisioning features to scale the cluster based on resource usage.
- C. Provision and configure each new environment and resources through the Google Cloud console.
- D. Develop custom Cloud Run functions to automate the provisioning of new Google Cloud environments. Incorporate security measures and pre-defined configurations within the functions to ensure consistency and protection.
- Answer

A

62.

For this question, refer to the EHR Healthcare case study. EHR Healthcare recently experienced unauthorized access of their Kubernetes-based applications due to compromised credentials, resulting in a security breach. You need to secure access between their on-premises systems and the Google Kubernetes Engine(GKE) environment, while maintaining regulatory compliance and integrating with EHR'S existing identity provider. What should you do?

- A. Establish a secure VPN tunnel between the on-premises network and the GKE VPC, and implement strict firewall rules to restrict access based on source IP addresses and ports.
- B. Create a service account with minimal permissions for each application running in GKE, and distribute the credentials securely to the on-premises systems.
- C. Configure Workload Identity Federation with Active Directory to enable authentication and authorization for applications across both Kubernetes environments.
- D. Deploy an OpenID Connect(OIDC) identity provider in GE, and integrate it with Active Directory for centralized authentication and fine-grained authorization controls.
- Answer

C

63.

For this question, refer to the EHR Healthcare case study. Due to the growing infrastructure administration costs, you want to migrate the legacy relational database systems to Google Cloud. You need to maintain the strict uptime requirements with limited in-house expertise. What should you do?

- A. Leverage AlloyDB to migrate the on-premises databases.
- B. Use the Database Migration Service(DMS) to migrate the MySQL databases to Cloud SQL.
- C. Backup and restore the databases to Cloud SQL instances with minimal changes.
- D. Set up database replication to Cloud SQL instances using native database tools, and incrementally switch applications to use the replicas.
- Answer

B

64.

For this question, refer to the Cymbal Retail case study. Cymbal wants to migrate their product catalog management processes to Google Cloud. You need to ensure a smooth migration with proper change management to minimize disruption and risks to the business. You want to follow Google-recommended practices to automate product catalog enrichment improve product discoverability, increase customer engagement, and minimize costs. What should you do?

- A Design a migration plan with a scheduled window to move all components at once. Perform extensive testing to ensure a successful migration.
- B. Design a migration plan to move all of Cymbal's data to Cloud Storage, and use Compute Engine for all business logic.
- C. Design a migration plan, starting with a pilot project focusing on a specific product category, and gradually expand to other categories.
- D. Design a migration plan to move all of Cymbal's data to Cloud Storage, and use Cloud Run functions for all business logic. Mark this item for later review.
- Answer

C

65.

For this question, refer to the Cymbal Retail case study. Cymbal wants you to connect their on-premises systems to Google Cloud while maintaining secure communication between their on-premises and cloud environments. You want to follow Google's recommended approach to ensure the most secure and manageable solution. What should you do?

- A. Configure a Cloud VPN gateway and establish a VPN tunnel. Configure firewall rules to restrict access to specific resources and services based on IP addresses and ports.
- B. Use a bastion host to provide secure access to Google Cloud resources from Cymbal's on-premises systems.
- C. Configure a static VPN connection using SSH tunnels to connect the on-premises systems to Google Cloud.
- D. Use Google Cloud's VPC peering to connect Cymbal's on-premises network to Google

Cloud.

- Answer

A

66.

For this question, refer to the Cymbal Retail case study. Cymbal plans to migrate their existing on-premises systems to Google Cloud and implement AI-powered virtual agents to handle customer interactions. You need to provision the compute resources that can scale for the AI-powered virtual agents. What should you do?

- A. Use Cloud SQL to store the customer data and product catalog.
- B. Create a single, large Compute Engine VM instance with a high CPU allocation.
- C. Deploy a Google Kubernetes Engine(GKE) cluster with autoscaling enabled.
- D. Configure Cloud Build to call AI Applications(formerly Vertex AI Agent Builder).
- Answer

C

67.

For this question, refer to the Cymbal Retail case study. Cymbal wants to migrate its diverse database environment to Google Cloud while ensuring high availability and performance for online customers. The company also wants to efficiently store and access large product images. These images typically stay in the catalog for more than 90 days and are accessed less and less frequently. You need to select the appropriate Google Cloud services for each database. You also need to design a storage solution for the product images that optimizes cost and performance. What should you do?

- A Migrate all databases to Spanner for consistency, and use Cloud Storage Standard for image storage.
- B. Migrate all databases to self-managed instances on Compute Engine, and use a

persistent disk for image storage.

- C. Migrate MySQL to Cloud SQL, SQL Server to Cloud SQL, Redis to Memorystore, and MongoDB to Firestore © Use Cloud Storage Standard for image storage, and move images to Cloud Storage Coldline storage when products become less popular.
- D. Migrate MySQL and SQL Server to Spanner, Redis to Memorystore, and MongoDB to Firestore. Use Cloud Storage Standard for image storage, and move images to Cloud Storage Nearline storage when products become less popular.
- Answer

C

68.

For this question, refer to the Cymbal Retail case study. Cymbal has a centralized project that supports large video files for Vertex AI model training. Standard storage costs have suddenly increased this month, and you need to determine why. What should you do?

- A. Investigate if the project owner moved from dual-region storage to region storage.
- B. Investigate if the project owner disabled a soft-delete policy on the bucket holding the video files.
- C. Investigate if the project owner moved from multi-region storage to region storage.
- D. Investigate if the project owner enabled a soft-delete policy on the bucket holding the video files.
- Answer

B

69.

For this question, refer to the Cymbal Retail case study. Cymbal's generative AI models require high- performance storage for temporary files generated during model training and inference. These files are ephemeral and frequently accessed and modified. You need to select a storage

solution that minimizes latency and cost and maximizes performance for generative AI workloads. What should you do?

- A. Use performance persistent disks.
- B. Use Local SSDs attached to the VMs running the generative AI models.
- C. Use Filestore to store temporary files.
- D. Use the Cloud Storage bucket in the same region as your virtual machines. Configure lifecycle policies to delete files after processing.
- Answer

B

70.

For this question, refer to the Cymbal Retail case study. Cymbal wants you to design a cloud-first data storage infrastructure for the product catalog modernization project. You want to ensure efficient data access and high availability for Cymbal's web application and virtual agents while minimizing operational costs. What should you do?

- A. Use Filestore for the structured product data, and Cloud Storage for product images.
- B. Use Spanner for the structured product data, and BigTable for product images.
- C. Use AlloyDB for structured product data, and Cloud Storage for product images.
- D. Use Cloud Storage for structured product data, and BigQuery for product images.
- Answer

C

71.

You are developing a retail application on Vertex AI that uses a large language model to give customers real-time product recommendations. The app uses text, code, images, audio, and video prompts to communicate with customers. You need to minimize the application's latency

to ensure a responsive and delightful user experience. What should you do?

- A. Avoid using system instructions to prevent overly long responses.
- B. Enable streaming responses from the model to process the output in real time.
- C. Increase the temperature parameter.
- D. Avoid using the max_output_tokens parameter to prevent overly long responses.
- Answer

B

72.

Your company is running AI/ML workloads on graphics processing units (GPUs) within Google Kubernetes Engine (GKE). These workloads are deployed across production and development projects, with the production environment being business critical and operating 24/7. You need to ensure that the GPU capacity is always available during scaling events, without being limited by overall GPU availability. You also need the ability to shift unused GPU capacity from production to development projects as needed. You want to follow Google-recommended best practices. What should you do?

- A. Create an owner project with shared reservations, and then configure the production project and development project as consumer projects to consume capacity.
- B. Create a reservation in the production project, and then transfer the reservation to the development project when needed.
- C. Create a shared reservation in the production project, and then add the development project as a consumer project to move capacity when needed.
- D. Create a reservation in the production project, and then create a separate reservation in the development project to secure capacity. Mark this item for later review
- Answer

A

73.

You are setting up a new regional environment for your business needs. You plan to migrate more than 2 TB of data from the us-central1 region to the asia-southeast1 region. The transfer should occur as quickly as possible. If the transfer is interrupted, the operation must resume where it left off and copy only new or modified objects. What should you do?

- A. Leverage Cloud Composer to orchestrate a data pipeline that copies objects from the source bucket to the destination bucket.
- B. Utilize Storage Transfer Service(STS) to migrate objects to the destination bucket in the target region. Once the transfer is complete, delete the original source bucket objects.
- C. Use the `gsutil storage op` command to copy the objects. Be careful to manage potential naming conflicts to ensure data integrity.
- D. Use `gsutil storage rsync-recursive` command to replicate the source bucket to a new target bucket in the us-west1 region.
- Answer

B

74.

You are responsible for the architecture of a critical ecommerce application. The application's web frontends run on Compute Engine instances deployed across multiple zones within a single region. The architecture must be highly available and self-healing. If a zone becomes unavailable, the system must automatically recreate the lost capacity in the available zones. To maximize efficiency, you need to ensure that all provisioned compute resources actively serve traffic during normal operations, and user traffic is automatically and seamlessly directed away from any failed instance or zone. What should you do?

- A. Create three separate zonal managed instance groups(MIGs), and use Cloud DNS with weighted routing to distribute traffic across the public IP address of an instance in each MIG.
- B. Configure a regional external Application Load Balancer with a primary regional managed instance group(MIG) as its backend service, and configure a second regional

MIG as a failover backend.

- C. Create a single regional managed instance group(MIG) configured to span all three zones. Place a regional external Application Load Balancer in front of the MIG, and configure its backend service with a health check.
- D. Create a single regional unmanaged instance group spanning all three zones. Place a regional external Application Load Balancer in front of the group, and configure its backend service with a health check.
- Answer

C

75.

You are designing the architecture for a stateless, public-facing, containerized web application. The marketing department uses this web application to run unpredictable flash sales that cause sudden, massive spikes in user traffic. The application is a critical revenue driver and must always be available and responsive to users. Your design must automatically handle high traffic fluctuations and allow for new application versions to be deployed with zero downtime and a safe, fast rollback capability. What should you do?

- A. Deploy the application to a regional Managed Instance Group(MIG). Configure a CPU-based autoscaling policy and use the MIG's rolling update feature for deployments.
- B. Deploy the application to a Google Kubernetes Engine(GKE) cluster. Use the Recreate deployment strategy to roll out new application versions.
- C. Deploy the application as a Cloud Run service. Configure request-based auto-scaling and use gradual rollouts by splitting traffic between revisions.
- D. Deploy the application to a single, large Compute Engine instance and use a startup script to pull and run the latest container image on boot.
- Answer

C

76.

You are designing the storage architecture for a financial analytics platform. The platform ingests and stores terabytes of transactional data daily, which is used for both real-time fraud detection and long-term historical analysis. Transaction data from the last 30 days must be accessible with very low latency for the fraud detection engine. Data older than 30 days is accessed infrequently for quarterly reports, where retrieval times of a few seconds are acceptable. All data must be retained for five years to meet compliance regulations. You need to design a solution as cost-effective as possible. What should you do?

- A. Store all transaction data in a Cloud Storage bucket using the Standard storage class for the entire five- year retention period.
- B. Ingest all data into BigQuery using time-partitioned tables, and rely on BigQuery's automatic long-term storage pricing for data older than 90 days.
- C. Configure a Cloud Storage bucket with an Object Lifecycle Management policy to transition data from the Standard class to the Archive class after 30 days.
- D. Configure a Cloud Storage bucket with an Object Lifecycle Management policy to transition data from the Standard class to the Coldline class after 30 days.
- Answer

D

77.

You are migrating a large, on-premises application to Google Cloud. The application consists of several interconnected virtual machines. You want to create a detailed migration plan to ensure a smooth migration with minimal effort. You need to understand the existing environment, identify dependencies, and estimate the total cost of ownership(TCO) in the cloud. What should you do?

- A. Use Config Connector to declare the desired state of your Google Cloud resources in Kubernetes-style manifests.
- B. Use the Google Cloud Migration Center to perform an automated discovery and assessment of the on-premises environment.

- C. Write a custom script to query the Sphere API for virtual machine information and then import it into BigQuery for analysis.
- D. Use the Google Cloud pricing calculator to input the specifications of your on-premises servers and receive a TCO estimate. Mark this item for later review
- Answer

B

78.

You are migrating a critical, on-premises, three-tier ecommerce application to Google Cloud. The application has tightly coupled dependencies between its web, application, and database tiers. The business has strict performance and availability requirements and is concerned about the risks associated with an "all-at-once" migration. You want to follow Google- recommended practices to create a migration plan that systematically addresses application dependencies and validates post-migration performance before the final production cutover. What should you do?

- A. Redesign the application to be cloud-native using microservices before migration to eliminate dependencies and simplify workload testing.
- B. Use dependency mapping to group application components into waves, create a phased migration plan, and incorporate performance testing for each wave.
- C. Plan a "lift and shift" migration for all application tiers to be executed simultaneously over a single weekend to minimize downtime.
- D. Prioritize the establishment of a high-bandwidth, low-latency Cloud Interconnect connection, and then plan to migrate the database servers first.
- Answer

B

79.

To improve governance and security, your organization has structured the Google Cloud

environment using folders for different business units. Each business unit folder has subfolders for development, staging, and production environments, which must comply with internal security controls: Production workloads must be protected from direct internet ingress by default unless explicitly tagged. The application must be accessible to customers over HTTPS. You need to design a scalable and enforceable model that blocks internet ingress traffic to the production folders while selectively allowing direct HTTPS traffic to the necessary virtual machines. You must also ensure that individual project teams cannot overwrite these controls once they are implemented for all current and future production projects. What should you do?

- A. At the organization root, apply a hierarchical firewall policy to deny all ingress except for HTTPS to tagged VMs.
- B. Mandate the application teams to deploy a Terraform module to create VPC firewall rules in each project that deny ingress and allow HTTPS.
- C. At each production folder, use an organization policy to block all external IPs and require teams to use external HTTPS load balancers.
- D. At each production folder, apply a hierarchical firewall policy to deny all ingress except for HTTPS to tagged VMs.
- Answer

D

80

You are designing the access control strategy for a new system that will store highly sensitive financial reports in a series of Cloud Storage buckets. Each bucket holds a specific type of sensitive data. Your company uses Google Workspace, and all employee roles are managed through Google Groups. The design must allow members of these groups to access only the reports they are authorized to see. The company's governance policy has two additional, critical requirements: All access must be directly tied to an employee's corporate identity for clear audit trails. The access model must be simple to operate and scalable. What should you do?

- A. Generate and distribute time-limited signed URLs to grant users temporary, per-object access.

- B Maintain fine-grained access control, and use default object ACLs to grant the required permissions to the Google Groups.
- C. Grant the IAM role roles/storage.objectAdmin to all authorized Google Groups to simplify permission management
- D. Enable uniform bucket-level access on the buckets, and grant predetined storage IAM roles to the Google Groups
- Answer

D

81.

A large enterprise is building a file processing pipeline on Google Cloud. The pipeline ingests the raw files, which are then accessed by a fleet of hundreds of Compute Engine virtual machines (VMs) in a single region for parallel processing. These VMs need concurrent, low-latency read and write access to a shared file system. You need to implement a storage solution that delivers the highest possible throughput to minimize the overall processing time. What should you do?

- A. Create a Filestore Zonal instance in each zone where the VMs are running. Mount the Filestore file share on each Compute Engine VM.
- B. Create a Filestore enterprise instance in the same region as the VMs. Mount the Filestore file share on each Compute Engine VM.
- C. Use a multi-region Cloud Storage bucket to store the files.
- D. Create a Regional Persistent Disk and attach it to all VMs in multi-writer mode.
- Answer

B

82.

Your company is building containerized applications as part of their CI/CD pipeline. To improve

the security and maintainability of the build process, you need to: Identify potential vulnerabilities within your container images. Generate verifiable metadata about the builds for auditing and compliance. Create a comprehensive inventory of your application's dependencies. What should you do?

- A. Use Cloud Build to build container images, and then trigger Artifact Analysis on images pushed to Artifact Registry
- B. Use Cloud Build to build container images, trigger Binary Authorization, and use cloud Asset Inventory for tracking and analysis.
- C. Use Cloud Build to build container images, trigger Binary Authorization, and use Security command Center for tracking and analysis.
- D. Use Cloud Build to build container images, push the images to Artifact Registry, and use Security Command Center for Tracking and analysis.
- Answer

A

83.

Your product team is building a critical, customer-facing application on Google Cloud. The development team wants to use Spanner for their database to take advantage of its horizontal scalability and low operational overhead. However, the FinOps team is concerned about the direct monthly cost of Spanner and proposed using a self-managed PostgreSQL database on Compute Engine VMs instead. You need to resolve this conflict and ensure the project moves forward with an architecturally sound database choice that balances technical requirements with financial constraints. What should you do?

- A. Provide the development team with a reference architecture for deploying a highly available PostgreSQL cluster on a regional managed instance group(MIG).
- B. Develop a total cost of ownership(TCO) analysis that includes operational overhead, and present it in a workshop to facilitate a decision.
- C. Cite the reliability and performance optimization pillars of the Google Cloud Well-Architected Framework to formally justify the use of Spanner.

- D. Suggest using Cloud SQL for PostgreSQL as a compromise to get a managed service at a lower cost than Spanner.
- Answer

D

84.

You are leading a large-scale Google Cloud implementation project. Your development teams need to deploy applications, but they lack deep expertise in cloud-native patterns. This leads to inconsistent and insecure deployments, such as GKE clusters without proper logging and overly permissive IAM roles. You need to implement a system that empowers these teams to deploy complete, pre-approved architectural patterns in a self-service manner. The system must ensure security and consistency by default while maintaining high developer velocity and avoiding manual approval bottlenecks. What should you do?

- A. Enforce a mandatory CI/CD pipeline using Cloud Build that requires every Terraform deployment to be reviewed and approved before the Terraform apply step can run.
- B. Create a central Cloud Monitoring dashboard and a set of custom alerts that fire whenever a non-compliant resource is detected.
- C. Use a combination of organization policies to restrict resource configurations and IAM recommender to periodically suggest improvements.
- D. Use Service Catalog to define, govern, and offer a portfolio of approved products. Include Terraform scripts into Service Catalog items.
- Answer

D

85.

You created a Cloud Monitoring dashboard for your dev environment to view real-time application metrics. The Ops team wants a similar dashboard to monitor the production environment. You want to minimize time and effort as you assist the Ops team with their

dashboard request. What should you do?

- A. Create a shared document and list every step you performed to create the dashboard. The Ops team can use the document to build their own dashboard.
- B. Share your dashboard and send the shared dashboard URL to the Ops team.
- C. Export your dashboard to a JSON file and share it with the Ops team.
- D. Grant viewer permission to the Ops team in your project so they can recreate the charts in their dashboard from the console
- Answer

C

86.

You are designing the observability strategy for a new microservices application running on Google Kubernetes Engine(GKE). The application consists of multiple services(e.g.frontend, orders, payments). During load testing, you observe an error in the frontend service's logs, but you cannot find the corresponding logs in the downstream services to investigate the root cause because the logs are not correlated. You need to implement a solution that allows you to follow a single user request across all microservices involved in the transaction. The solution must not require developers to manually add correlation logic to their application code. What should you do?

- A. Require developers to generate a unique correlation-id at the frontend, and manually add it as a field to every log message in all services.
- B. Configure all containers to write logs to STDOUT/STDERR, and then filter logs by pod name and timestamp in the Logs Explorer.
- C. Create custom metrics in Cloud Monitoring for error counts in each service and correlate incident spikes using a shared dashboard.
- D. Implement Cloud Trace by ensuring the traceparent header is propagated between microservice calls to link logs to a single trace.
- Answer

D

87.

Your organization has a significant amount of log data stored in Cloud Logging. The data engineering team is accustomed to using SQL for analysis and wants the ability to create insightful dashboards for visualizing log trends and patterns. You want to follow the recommendations of the Google Cloud Well-Architected Framework to provide a solution for the data engineering team. What should you do?

- A. Enable log analytics and run queries in Cloud Monitoring. Visualize the data using Vertex AI workbench.
- B. Create a log sink, and export the data to a storage bucket. Create an external table in BigQuery for the data in the bucket. Run queries and visualize the data with Cloud Monitoring dashboards.
- C. Enable log analytics and run queries in the linked log dataset in Big Query. Visualize the data with Looker Studio dashboard
- D. Create a log sink, and export the data to BigQuery using Pub/Sub. Run queries and visualize the data with Cloud Monitoring dashboards.
- Answer

C

88.

For this question, refer to the EHR Healthcare case study. You are responsible for designing the Google Cloud network architecture for Google Kubernetes Engine. You want to follow Google best practices. Considering the EHR Healthcare business and technical requirements, what should you do to reduce the attack surface?

- A. Use a public cluster with firewall rules and Virtual Private Cloud(VPC) routes.
- B. Use a private cluster with a public endpoint with authorized networks configured.

- C. Use a private cluster with a private endpoint with authorized networks configured.
- D. Use a public cluster with authorized networks enabled and firewall rules.
- Answer

C

89.

For this question, refer to the EHR Healthcare case study. You are responsible for ensuring that EHR's use of Google Cloud will pass an upcoming privacy compliance audit. What should you do?

- A. Verify EHR's product usage against the list of compliant products on the Google Cloud compliance page.
- B. Use GKE private clusters for all Kubernetes workloads.
- C. Use Firebase Authentication for EHR's user-facing applications.
- D. Implement Prometheus to detect and prevent security breaches on EHR's web-based applications.
- Answer

A

90.

For this question, refer to the EHR Healthcare case study. You are a developer on the EHR customer portal team. The application servers have increased load, and the application is having timeout errors. You recently incorporated Pub/Sub into the application architecture, and the application is not logging any Pub/Sub publishing errors. You want to improve publishing latency. What should you do?

- A. Increase the Pub/Sub Total Timeout retry value.
- B. Create a backup Pub/Sub message queue.

- C. Move from a Pub/Sub subscriber pull model to a push model.
- D. Turn off Pub/Sub message batching.
- Answer

D

91.

For this question, refer to the EHR Healthcare case study. You need to define the technical architecture for hybrid connectivity between EHR's on-premises systems and Google Cloud. You want to follow Google- recommended practices for production-level applications. Considering the EHR Healthcare business and technical requirements, what should you do?

- A. Configure two Partner Interconnect connections in one metro(City), and make sure the Interconnect connections are placed in different metro zones.
- B. Configure two VPN connections from on-premises to Google Cloud, and make sure the VPN devices on-premises are in separate racks.
- C. Configure two Dedicated Interconnect connections in one metro(City) and two connections in another metro,and make sure the Interconnect connections are placed in different metro zones.
- D. Configure Direct Peering between EHR Healthcare and Google Cloud, and make sure you are peering at least two Google locations.
- Answer

C

92.

For this question, refer to the EHR Healthcare case study. You need to define the technical architecture for securely deploying workloads to Google Cloud. You also need to ensure that only verified containers are deployed using Google Cloud services. What should you do?

- A Enable Binary Authorization on GKE, and sign containers as part of a CI/CD pipeline.

- B. Configure Artifact Registry to only allow trusted service accounts to create and deploy containers from the registry.
- C. Configure Artifact Registry to use vulnerability scanning to confirm that there are no vulnerabilities before deploying the workload.
- D. Configure Jenkins to utilize Kritis to cryptographically sign a container as part of a CI/CD pipeline.
- Answer

A

93.

The operations team in your company wants to save Cloud VPN log events for one year. You need to configure the cloud infrastructure to save the logs. What should you do?

- A. Enable the Compute Engine API, and then enable logging on the firewall rules that match the traffic you want to save.
- B. Set up a Cloud Logging Dashboard titled Cloud VPN Logs, and then add a chart that queries for the VPN metrics over a one-year time period.
- C. Set up a filter in Cloud Logging and a Cloud Storage bucket as an export target for the logs you want to save.
- D. Set up a filter in Cloud Logging and a topic in Pub/Sub to publish the logs.
- Answer

C